

Predicting bird species from their songs using machine learning

Theme 09 - Introduction Machine Learning
Research paper



Figure 1: A common redpoll (*Acanthis Flammea*)

Niek R. Scholten (388602)
Bio-Informatics BFV3
Institute of Life Science & Technology
Hanze University of Applied Sciences
Dave Langers (LADR) & Bart Barnard (BABA)
November 13, 2022
From: <https://github.com/niek265/Thema9-Praktijkopdracht>

Predicting bird species from their songs using machine learning

Theme 09 - Introduction Machine Learning

Niek R. Scholten (388602)

Bio-Informatics BFV3

Institute of Life Science & Technology

Hanze University of Applied Sciences

Dave Langers (LADR) & Bart Barnard (BABA)

November 13, 2022

From: <https://github.com/niek265/Thema9-Praktijkopdracht>

Abstract

The complexity and the consistency of birds' songs have long been marveled over. And they always say "consistency is key", which is also true for the models in this project. All the data that was used follows nice conventions and is perfect for machine learning. This study aims to find out how well new songs of those same bird species can be classified using a model trained on the original data using machine learning.

Summary

Birds' songs can be complex and hard to distinguish by ear, but these songs can be easily transcribed to numeric data which is easily understood by computers. The goal of this research is to create an easy to run program that can classify new instances of this numeric convention. Python was used in pre-processing to make sure all data was structured correctly and ready for visualisation. Weka was used after this to explore possible models for the dataset, the best model was chosen and exported to be used in a Java program. The wrapper has been successfully built and can classify new instances with a high accuracy.

List of Abbreviations

CSV Comma-separated values

STFT Short-time Fourier transform

Contents

Abstract	i
Summary	ii
List of Abbreviations	iii
1 Introduction	1
1.1 Purpose	1
1.2 Theory	1
2 Materials & Methods	2
2.1 Materials	2
2.2 Methods	3
3 Results	6
3.1 Exploration	6
3.2 Model performance	9
4 Conclusion	11
5 Discussion	11
6 References	12

1 Introduction

1.1 Purpose

The wonders of birds' songs can best be experienced in person, but to accurately predict the species can be difficult. This research aims to provide an easy way to predict the species of a bird song.

1.2 Theory

1.2.1 Sound theory

The provided data has 2 different kinds features; chroma features & spectral centroids. These features are calculated on 13 intervals from a pre-defined random "window" of a bird's song. There are 12 chroma features and 1 spectral centroid per interval. So in the end each sample of a song is converted to 168 numeric data points (13 features * 13 intervals).

1.2.1.1 Chroma features capture the melodic characteristics of music, they capture the intensity of a certain tone at any given time. Each of the tone heights in an octave fits into a set of chroma features: {C, C#, D, D#, E, F, F#, G, G#, A, A#, B}. These features are calculated using STFT's (Short-time Fourier transform).

$$\hat{f}(\xi) = \begin{cases} \int_{-\infty}^{\infty} f(x) e^{-i2\pi\xi x} dx, & \xi \geq 0 \\ \hat{f}^*(|\xi|) & \xi < 0 \end{cases},$$

(1.2.1.1 A Fourier transform, x represents time. [1])

A STFT uses equation 1.2.1.1 on every interval of an audio clip to generate the chroma features.

1.2.1.2 Spectral centroid is described as the center of mass of a spectrum. It is also known as the "brightness" or timbre of sound. It is also calculated using a Fourier transform 1.2.1.1 with the magnitude used as weights.

$$\text{Centroid} = \frac{\sum_{n=0}^{N-1} f(n)x(n)}{\sum_{n=0}^{N-1} x(n)}$$

(1.2.1.2 The formula for calculating the centroid, $x(n)$ represents magnitude. [2])

1.2.2 Algorithm theory

The algorithm used for this research is called a random forest. It combines regular decision trees with bagging to create a "forest" of randomized decision trees. Regular decision trees are very prone to overfitting (matching the original data too much) the deeper they get, increasing irregularity. The reason random forest is so powerful is that it averages multiple decision trees to create a model that's not as biased as regular decision trees and less irregular.

1.2.2.1 Bagging splits up a data set and trains the model on these splits, only to recombine them later on. This greatly reduces overfitting and increases variance.

1.2.2.2 Decision trees create leaves that represent the class labels and branches that represent decision paths. A new instance can then be entered into it, so it will end at one of these leaves in accordance to its values.

2 Materials & Methods

The dataset contains converted data from bird song samples. The datapoints contain information about the intensity of certain tones in these samples. The CSV data was supplied by Edoardo Ferrante on Kaggle. This data was created using the Librosa package for Python. Librosa outputs the intensity of a certain tone at different time intervals from the provided sound file. There is a Python script available from the author that explains how the data was transformed from sound files to numeric data points, but this script has a specific shortcoming that will be explained later.

2.1 Materials

Several pieces of software were used over the course of this research, from packages to programs. The data was supplied by Edoardo Ferrante on Kaggle and contains CSV (comma-separated values) files for both training and testing.

2.1.1 Programs

Software	Package	Version
R	Dict	0.1.0
	dplyr	1.0.9
	ggbiplot2	0.55
	ggplot2	3.3.6
	ggpubr	0.4.0
	knitr	1.40
	plyr	1.8.7
	RWeka	0.4-44
	scales	1.2.1
		3.10
Python	librosa	0.9.2
	matplotlib	3.5.3
	numpy	1.23.3
	pandas	1.4.4
		3.8.6
Weka	wekaDeepLearning4j	1.7.2
		3.8.6
Java	openJDK17	17.0.4
Gradle		7.4
	weka-stable	3.8.6

Table 1: Programs and versions

The data consists of multiple iterations of the original CSV files. A metadata file is also included as to give some extra context to the source of the audio files.

2.2 Methods

Pre-processing and model training are the most intense operations performed in this research. There was also a bit of light data exploration which is contained within the log.

2.2.1 Data Pre-Processing

First let's have a look at how Librosa normally outputs the chromogram data:

	0	1	2	3	4	5	6
chromogram_0	0.68661	0.67378	0.65758	0.66149	0.68533	0.72239	0.76395
chromogram_1	0.91368	0.88148	0.85024	0.82476	0.82282	0.83024	0.83908
chromogram_2	0.98221	0.97060	0.95834	0.94729	0.94785	0.95189	0.95408
chromogram_3	1.00000	1.00000	1.00000	1.00000	1.00000	1.00000	1.00000
chromogram_4	0.96223	0.95790	0.95436	0.95403	0.94404	0.93285	0.91552
chromogram_5	0.92098	0.89960	0.88007	0.86290	0.84262	0.82280	0.80381
chromogram_6	0.87591	0.85544	0.83320	0.80870	0.78241	0.75819	0.73456
chromogram_7	0.79397	0.79418	0.79848	0.80535	0.80741	0.80750	0.80913
chromogram_8	0.62856	0.64859	0.68255	0.72441	0.76565	0.80078	0.82178
chromogram_9	0.41881	0.41914	0.44003	0.48525	0.54615	0.61608	0.68581
chromogram_10	0.36895	0.38299	0.40678	0.43114	0.45138	0.47237	0.52333
chromogram_11	0.33855	0.34676	0.35929	0.37378	0.38602	0.39931	0.42756

Table 2: Librosa output data

As you can see, the output is a neat array containing the values of 12 different tones at different time intervals. This data is sorted and can be read by Librosa.

Now onto the issue; the provided dataset contains this data in a stacked order, so each sample only takes up one row. This is a good idea, but due to sorting by alphabetical order the original order is lost. The order is important because we are working with data over time. This is not a problem if the trained model is only used on the provided test data, but we want the trained model to work in as many situations as possible.

Here is a look at the provided data:

id	chromogram_0_0	chromogram_0_1	chromogram_0_10	chromogram_0_11
0	0.997943662321316	0.832392210770135	0.7653861625931	0.70427464132375
1	0.996254885931866	0.839119599044146	0.760416790506312	0.705141765139875
2	0.970810156116343	0.823539694937237	0.759508104372184	0.709057883677716
3	1	0.855558393364941	0.752038009313116	0.710976936190937
4	1	0.884304523555434	0.741884532311754	0.714775207828629
5	0.971867873978603	0.824311712155432	0.755293860709407	0.71448132195049
6	1	0.835499361583387	0.751917158063063	0.717361992854453
7	0.978929855885584	0.827216718543843	0.751072631712318	0.718400862681119
8	1	0.895339720206626	0.733409813021178	0.722747412968086
9	0.967651828343747	0.823697857901917	0.746005680687241	0.721194823494439
10	0.993699774531599	0.847257121555946	0.734368883301346	0.726420069139032
11	0.00947350497274455	0.00699383738737368	0.372026644035831	0.0516494292032762
12	0.00982270123521504	0.00712337798131429	0.371129653847745	0.051631441504244

Table 3: Provided data

Each row contains a stack of chromogram data in a non-sequential order. The end of the array also contains the species of the corresponding bird and some spectral centroid data.



Figure 2: Comparison of the data order

Figure 2 shows a comparison of the order of provided data and the ideal order the data should be sorted in. A Python script was used to re-order the data into this format.

Data exploration was not a big part of this research, the data was fairly straightforward after this initial processing. Some exploration was done to look at differences between the species, this is further fleshed out in the results.

While exploring the data, all instances were exported to arff files using the RWeka package. In Weka these were subsequently altered so the class labels were nominal values and placed in the last column. An extra test data set was also created with all of its instances randomized, so there was no bias to the order.

2.2.2 Model training

To create a usable model in Weka, a lot of algorithms were tested for efficiency. Namely:

Algorithm
ZeroR
OneR
NaiveBayes
Sequential minimal optimization
SimpleLogistic
IBk
J48 Tree
RandomForest
wekaDeeplearning4j

Table 4: Tested algorithms

The option of wekaDeeplearning4j was explored, but due to the difficult nature to set up and lackluster performance compared to cheaper algorithms it was not explored further.

All models were trained on the train dataset using 10-fold cross-validation with varying results. The random forest model file was exported for later use in the wrapper.

2.2.3 Wrapper creation

The finalized wrapper was built using Gradle and the weka-stable package for Java. The aforementioned model file containing the trained random forest model is incorporated into the jar file and shipped with the program. It was tested on the unlabeled test data set and validated afterwards using the correct labels.

3 Results

The results of this research come in 2 sections, exploration and model performance. The exploration section contains some results of the exploratory data analysis along with figures to sketch an idea of the data that was used. The section on model performance contains detailed information about the model that was chosen as final classifier.

3.1 Exploration

Chroma comparison for 2 birdsongs

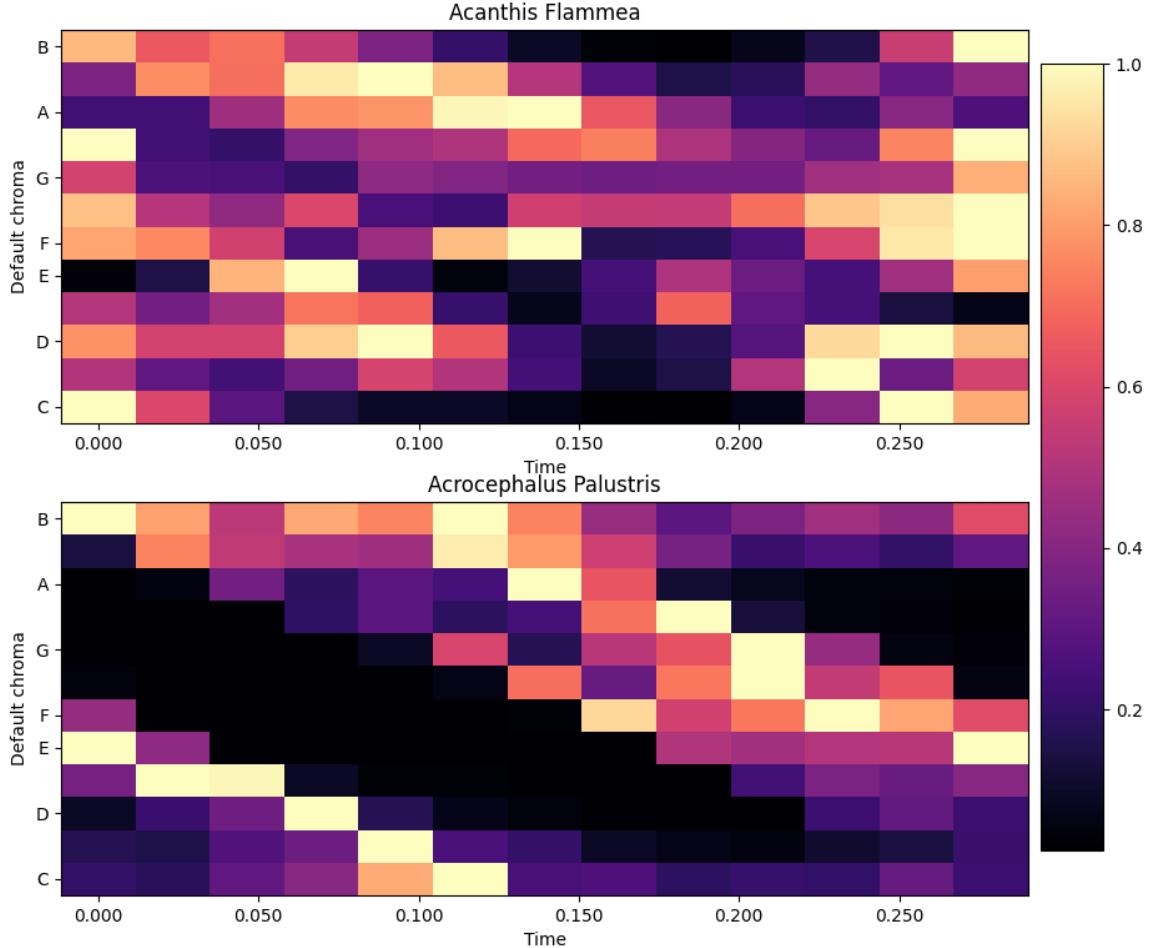


Figure 3: Chroma signature comparison for 2 fragments of different bird-species songs

Figure 3 shows the first result of the data exploration. There is a clear difference in signature between the two species. This difference can be explored further by analyzing one tone at the time.

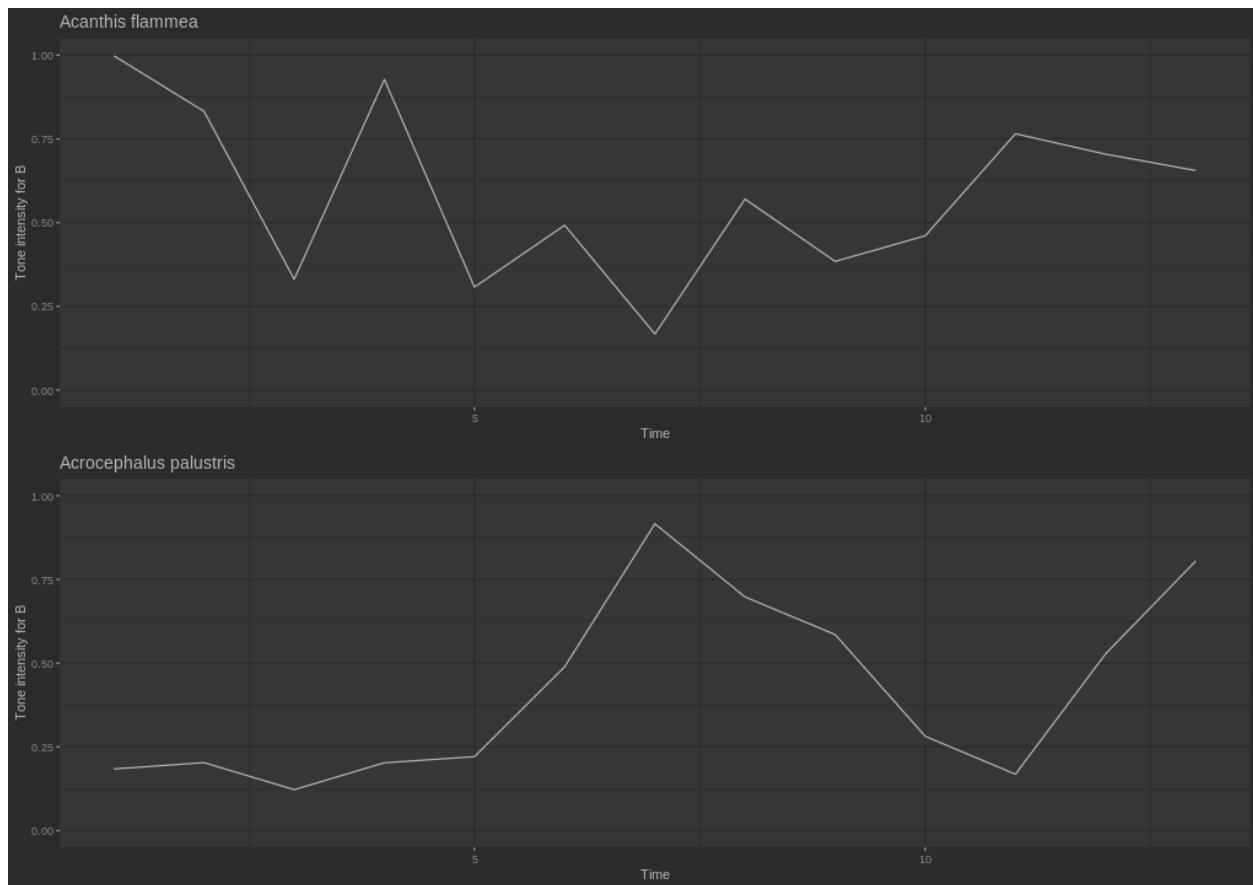


Figure 4: Tone intensity comparison at tone B for 2 fragments of different bird-species songs

Figure 4 also shows a clear difference in tone intensity over the duration of the sound fragment. Tone B was chosen as an example, but multiple tones exhibit this behaviour as seen in Figure 2.

Chroma comparison for 2 birdsongs

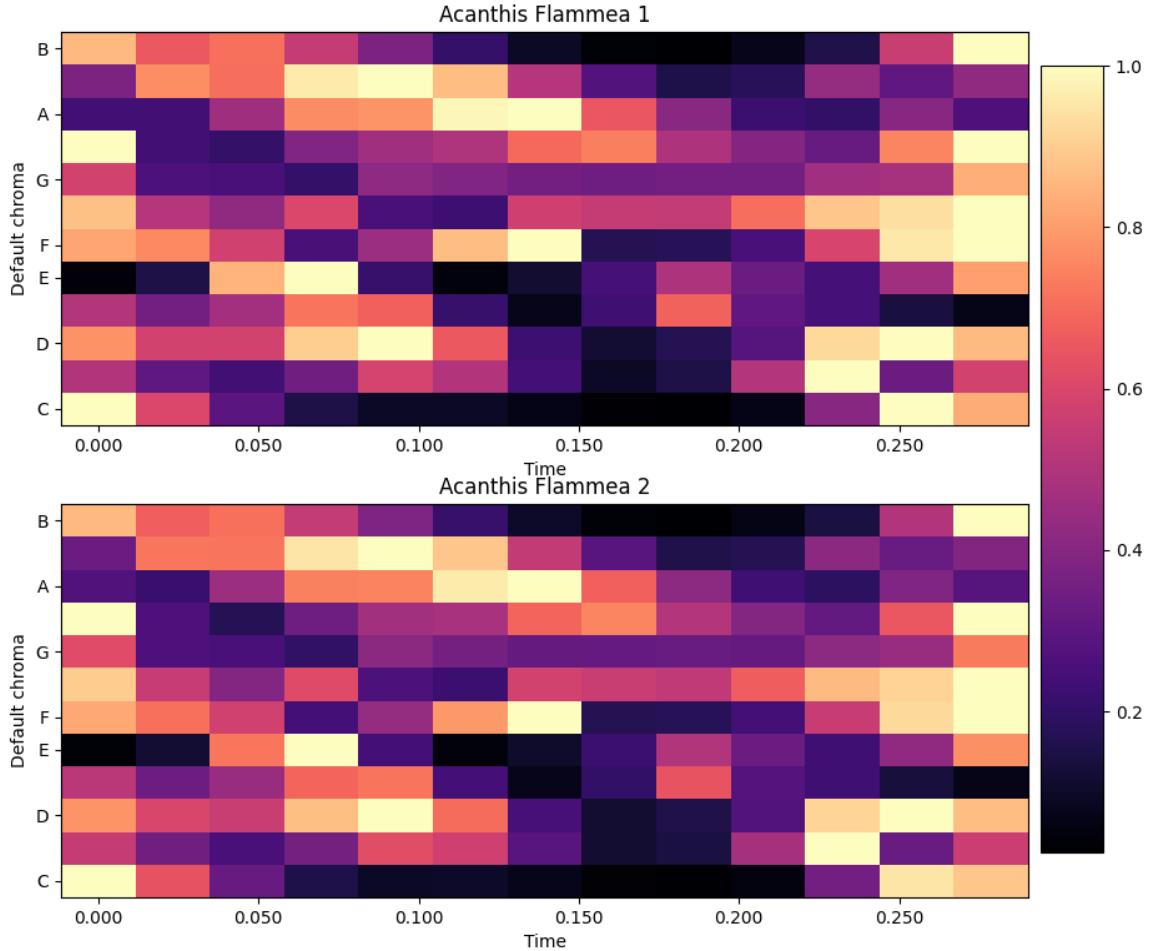


Figure 5: Chroma signature comparison for 2 fragments of the same bird-species songs

Lastly, figure 5 shows the signature comparison of 2 bird songs of the same species. These signatures look alike, but there are subtle differences to be seen.

3.2 Model performance

As mentioned in previous chapters, the random forest classifier is the best one for this specific use case. It was chosen by evaluating different models using the Weka experimenter. Pinning these models against each other yielded the following results:

Algorithm	% correct
ZeroR	1.14
OneR	52.45
NaiveBayes	92.12
Sequential minimal optimization	98.06
SimpleLogistic	98.02
IBk	98.13
J48 Tree	93.40
RandomForest	98.55

Table 5: Tested algorithms and their performance

As seen above, the performance of the random forest is almost matched in a few cases, but with some tweaks a % correct score of 98.8632 % was achieved. Here are the parameters used to achieve this result:

```
weka.classifiers.trees.RandomForest -- -P 100 -I 110 -num-slots 12 -K 0 -M 1.0 -V 0.001 -S 1
```

Another upside of the random forest classifier is its speed, it is able to work in parallel and increase performance linearly. Other models take a lot of time to train and test using cross validation.

OneR also scores pretty high, indicating that one label can be enough to correctly classify most of the data. In this case the spectral centroids hold a lot of information, as this is what OneR uses as label. If this specific spectral centroid is removed from the dataset, the performance doesn't drop that much since it will use another spectral centroid. However, if all spectral centroids are removed, OneR's performance falls to 30%. This indicates that spectral centroids hold a lot of information about the birds' species.

Because the train and test set are of differing sizes, they are incompatible. Therefore, the model has to be passed through `weka.classifiers.misc.InputMappedClassifier` to be able to test on the test set. This is also incorporated in the Java wrapper.

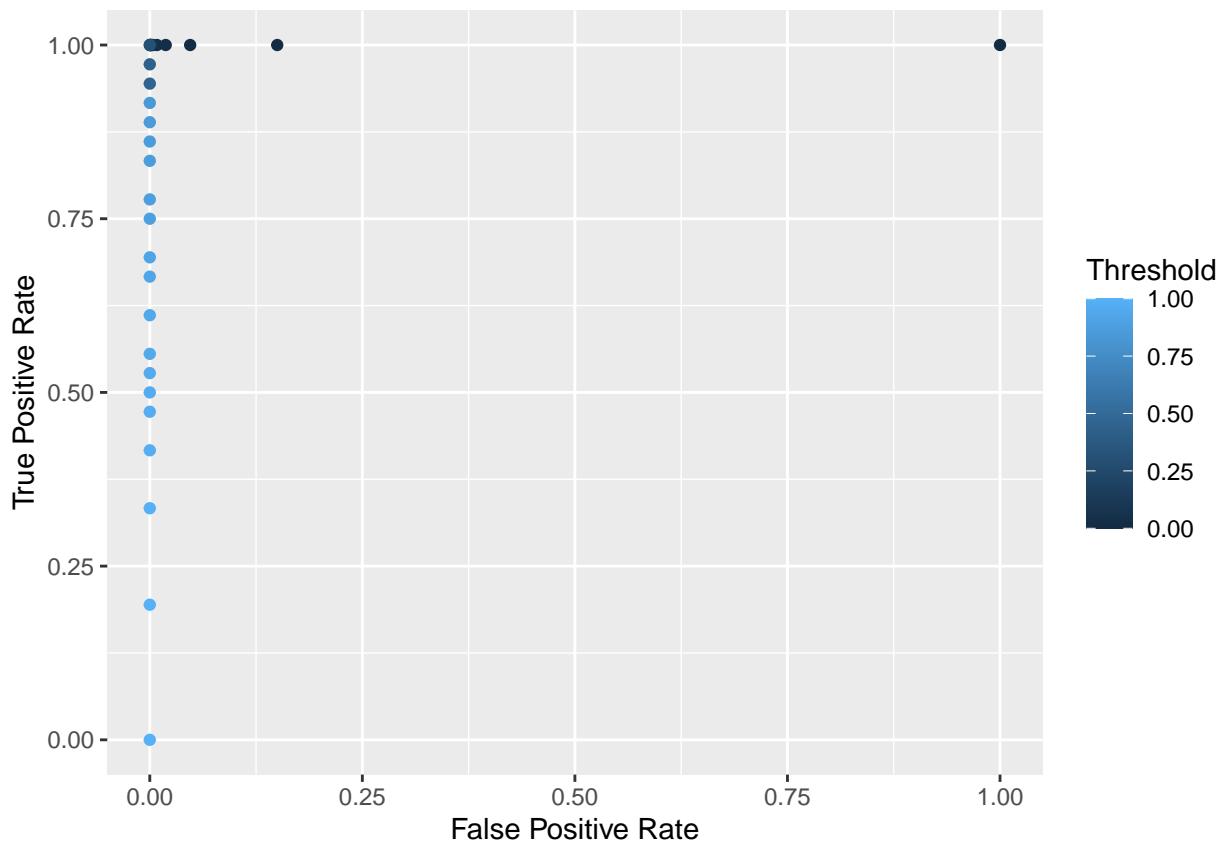


Figure 6: ROC Curve of *Acanthis Flammea*

The figure above shows the ROC Curve for *Acanthis Flammea*, the model is able to classify all instances of the test set perfectly. The training set is balanced, containing 20 samples of each species, this attributes to the great results of this model.

4 Conclusion

Concluding from the data, it seems there is a high probability of success when training a machine learning model on this data. The signatures of the same species align with each other nicely and the signatures of other species are different enough to be separated. Although the signatures of the same species look alike, there are subtle differences in the intensity. We can also conclude from this that birds of the same species sing the same songs.

And from the model creation we can conclude that this was an absolute success. It can predict new instances with an extremely high accuracy and few incorrectly classified instances. The balance between high accuracy and overfitting is just right, exceeding expectations.

5 Discussion

The model is not perfect, if it was it would have achieved perfect scores over the entire test set. A randomized train set could be created by merging the current train and test sets to form one large data set. This could potentially increase the performance even more as there is more information variance available, but the current model already achieves high performance with just 20 entries of every species.

The model could also be expanded to work on longer fragments of songs or more bird species altogether.

6 References

1. Chroma Feature Extraction, 2019 Ayush Kumar Shah: ResearchGate 330796993
2. Brain Seizure Detection and Classification Using EEG Signals, 2022 Varsha K.Harpale: ScienceDirect B978-0-32-391120-7.00008-6